



Hive-HBase Integration

IMPORTANT

Copyright Infringement and Illegal Content Sharing Notice

All course content designs, video, audio, text, graphics, logos, images are Copyright© and are protected by India and international copyright laws. All rights reserved.

Permission to download the contents (wherever applicable) for the sole purpose of individual reading and preparing yourself to crack the interview only. Any other use of study materials – including reproduction, modification, distribution, republishing, transmission, display – without the prior written permission of Author is strictly prohibited.

Trendytech Insights legal team, along with thousands of our students, actively searches the Internet for copyright infringements. Violators subject to prosecution.



What are we trying to do?

We will see how to create a table which we can access from hive as well as hbase.

Remember: we want to access this table through hive if we want to do any processing or analysis.

We want to access this table through hbase for quick searching or transactional activities.

Hive-hbase integration

Step 1: Make sure file kv1.txt file is placed inside Downloads folder in cloudera VM (the delimiter in this file is CTRL-A character which is default in delimiter in hive)

```
[cloudera@quickstart Downloads]$ head kv1.txt
238val 238
86val 86
311val 311
27val 27
165val 165
409val 409
255val 255
278val 278
98val 98
484val 484
[cloudera@quickstart Downloads]$
```



Hive-hbase integration

Step 2: Create a table in hive

CREATE TABLE pokes (foo INT, bar STRING);

```
hive> CREATE TABLE pokes (foo INT, bar STRING);  
OK  
Time taken: 0.929 seconds  
hive> █
```



Hive-hbase integration

Step 3: Load the data in hive table

LOAD DATA LOCAL INPATH '/home/cloudera/Downloads/kv1.txt'
OVERWRITE INTO TABLE pokes;

```
[cloudera@quickstart ~]$ hive

Logging initialized using configuration in file:/etc/hive/conf.dist/hive-log4j.properties
WARNING: Hive CLI is deprecated and migration to Beeline is recommended.
hive>
hive> CREATE TABLE pokes (foo INT, bar STRING);
OK
Time taken: 0.929 seconds
hive> LOAD DATA LOCAL INPATH '/home/cloudera/Downloads/kv1.txt' OVERWRITE INTO TABLE pokes;
Loading data to table default.pokes
Table default.pokes stats: [numFiles=1, numRows=0, totalSize=5812, rawDataSize=0]
OK
Time taken: 0.648 seconds
```




Hive-hbase integration

Step 4: verify the data in hive table

```
SELECT * FROM pokes WHERE foo = 98;
```

```
hive> SELECT * FROM pokes WHERE foo = 98;  
OK  
98      val_98  
98      val_98  
Time taken: 0.51 seconds, Fetched: 2 row(s)  
hive> █
```

The output of the SELECT command displays two identical rows because there are two identical rows in the Hive pokes table with a key of 98.

Note: This is a good illustration of the concept that Hive tables can have multiple identical keys. As we will see shortly, HBase tables cannot have multiple identical keys, only unique keys.



Hive-hbase integration

step 5: create a Hive-HBase table

```
CREATE TABLE hbase_table_1(key int, value string) STORED BY  
'org.apache.hadoop.hive.hbase.HBaseStorageHandler' WITH  
SERDEPROPERTIES ("hbase.columns.mapping" = ":key,cf1:val")  
TBLPROPERTIES ("hbase.table.name" = "xyz");
```

Note: The TBLPROPERTIES command is not required, but those new to Hive-HBase integration may find it easier to understand what's going on if Hive and HBase use different names for the same table.

In this example, Hive will recognize this table as "hbase_table_1" and HBase will recognize this table as "xyz".

Hive-hbase integration

```
hive> CREATE TABLE hbase_table_1(key int, value string) STORED BY 'org.apache.hadoop.hive.hbase.HBase
StorageHandler' WITH SERDEPROPERTIES ("hbase.columns.mapping" = ":key,cf1:val") TBLPROPERTIES ("hbase
.table.name" = "xyz");
OK
Time taken: 2.685 seconds
```

Check the tables in hive using show tables command

```
hive> show tables;
OK
hbase_table_1
pokes
Time taken: 0.068 seconds, Fetched: 2 row(s)
hive> █
```

We can see **hbase_table_1**



Hive-hbase integration

Step 6: in hbase check the list of table

list

```
hbase(main):003:0* list
TABLE
census
xyz
2 row(s) in 0.0090 seconds
=> ["census", "xyz"]
```

Note: HBase recognizes the Hive-HBase table named xyz. This is the same table known to Hive as hbase_table_1.

Hive-hbase integration

step 7: From the Hive prompt, insert data from the Hive table pokes into the Hive-HBase table hbase_table_1

INSERT OVERWRITE TABLE hbase_table_1 SELECT * FROM pokes WHERE foo=98;

```
hive> INSERT OVERWRITE TABLE hbase_table_1 SELECT * FROM pokes WHERE foo=98;
Query ID = cloudera_20200526125454_72ec09ab-59c7-4bcc-8582-fce8ee50b71f
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks is set to 0 since there's no reduce operator
Starting Job = job_1590135721705_0001, Tracking URL = http://quickstart.cloudera:8088/proxy/application_1590135721705_0001/
Kill Command = /usr/lib/hadoop/bin/hadoop job -kill job_1590135721705_0001
Hadoop job information for Stage-0: number of mappers: 1; number of reducers: 0
2020-05-26 12:55:06,509 Stage-0 map = 0%, reduce = 0%
2020-05-26 12:55:18,920 Stage-0 map = 100%, reduce = 0%, Cumulative CPU 2.68 sec
MapReduce Total cumulative CPU time: 2 seconds 680 msec
Ended Job = job_1590135721705_0001
MapReduce Jobs Launched:
Stage-Stage-0: Map: 1 Cumulative CPU: 2.68 sec HDFS Read: 17508 HDFS Write: 0 SUCCESS
Total MapReduce CPU Time Spent: 2 seconds 680 msec
OK
Time taken: 31.064 seconds
hive> █
```



Hive-hbase integration

step 8: In hive Query hbase_table_1 to see the data we have inserted into the Hive-HBase table:

```
SELECT * FROM hbase_table_1;
```

```
hive> SELECT * FROM hbase_table_1;  
OK  
98      val_98  
Time taken: 0.149 seconds, Fetched: 1 row(s)  
hive> █
```

Note: Note: Even though we loaded two rows from the Hive pokes table that had the same key of 98, only one row was actually inserted into hbase_table_1. This is because hbase_table_1 is an HBASE table, and although Hive tables support duplicate keys, HBase tables only support unique keys. HBase tables arbitrarily retain only one key, and will silently discard all the data associated with duplicate keys.

Hive-hbase integration

Step 9: let's try to see the same data using hbase shell

scan 'xyz'

```
hbase(main):005:0* scan 'xyz'
ROW
 98
1 row(s) in 0.2510 seconds
COLUMN+CELL
column=cf1:val, timestamp=1590522917727, value=val_98
hbase(main):006:0> █
```



We have learnt Hive-Hbase Integration

Happy Learning!!!



5 Star Google Rated
Big Data Course

LEARN FROM THE EXPERT



9108179578

Call for more details



Follow US

Trainer Mr. Sumit Mittal

Phone 9108179578

Email trendytech.sumit@gmail.com

Website <https://trendytech.in/courses/big-data-online-training/>

LinkedIn <https://www.linkedin.com/in/bigdatabysumit/>

Twitter @BigdataBySumit

Instagram bigdatabysumit

Facebook <https://www.facebook.com/trendytech.in/>

Youtube TrendyTech